



PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Preprocessing Data for Machine Learning

Sarah Guido
Senior Data Scientist

What is data preprocessing?

- Beyond cleaning and exploratory data analysis
- Prepping data for modeling
- Modeling in Python requires numerical input

Refresher on Pandas basics

```
In [1]: import pandas as pd
```

```
In [2]: hiking = pd.read_json("datasets/hiking.json")
```

```
In [3]: print(hiking.head())
```

	Accessible	Difficulty	Length	Limited_Access
0	Y	None	0.8 miles	N
1	N	Easy	1.0 mile	N
2	N	Easy	0.75 miles	N
3	N	Easy	0.5 miles	N
4	N	Easy	0.5 miles	N

Refresher on Pandas basics

```
In [4]: print(hiking.columns)

Index(['Accessible', 'Difficulty',
       'Length', 'Limited_Access',
       'Location', 'Name',
       'Other_Details', 'Park_Name',
       'Prop_ID', 'lat', 'lon'],
      dtype='object')
```

```
In [5]: print(hiking.dtypes)
```

Accessible	object
Difficulty	object
Length	object
Limited_Access	object
Location	object
Name	object
Other_Details	object
Park_Name	object
Prop_ID	object
lat	float64
lon	float64
dtype:	object

Refresher on Pandas basics

```
In [6]: print(wine.describe())
```

	Type	Alcohol	Malic acid	Ash	Alcalinity of ash
count	178.000000	178.000000	178.000000	178.000000	178.000000
mean	1.938202	13.000618	2.336348	2.366517	19.494944
std	0.775035	0.811827	1.117146	0.274344	3.339564
min	1.000000	11.030000	0.740000	1.360000	10.600000
25%	1.000000	12.362500	1.602500	2.210000	17.200000
50%	2.000000	13.050000	1.865000	2.360000	19.500000
75%	3.000000	13.677500	3.082500	2.557500	21.500000
max	3.000000	14.830000	5.800000	3.230000	30.000000

Removing missing data

```
In [7]: print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
In [8]: print(df.dropna())
```

	A	B	C
1	4.0	7.0	3.0
4	5.0	9.0	7.0

Removing missing data

```
In [9]: print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
In [10]: print(df.drop([1, 2, 3]))
```

	A	B	C
0	1.0	NaN	2.0
4	5.0	9.0	7.0

Removing missing data

```
In [11]: print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
In [12]: print(df.drop("A", axis=1))
```

	B	C
0	NaN	2.0
1	7.0	3.0
2	NaN	NaN
3	7.0	NaN
4	9.0	7.0

Removing missing data

```
In [13]: print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
In [14]: print(df[df["B"] == 7])
```

	A	B	C
1	4.0	7.0	3.0
3	NaN	7.0	NaN

Removing missing data

```
In [15]: print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
In [16]: print(df["B"].isnull().sum())
```

```
2
```

```
In [17]: print(df[df["B"].notnull()])
```

	A	B	C
1	4.0	7.0	3.0
3	NaN	7.0	NaN
4	5.0	9.0	7.0



PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Let's practice!



PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Working With Data Types

Why are types important?

```
In [1]: print(volunteer.dtypes)

opportunity_id      int64
content_id          int64
vol_requests        int64
event_time          int64
title               object
hits                int64
summary             object
is_priority          object
category_id         float64
...
```

- object: string/mixed types
- int64: integer
- float64: float

Converting column types

```
In [2]: print(df)
```

	A	B	C
0	1	string	1.0
1	2	string2	2.0
2	3	string3	3.0

```
In [3]: print(df.dtypes)
```

A	int64
B	object
C	object

dtype: object

Converting column types

```
In [4]: print(df)

   A      B      C
0  1  string  1.0
1  2 string2  2.0
2  3 string3  3.0

In [5]: df["C"] = df["C"].astype("float")
In [6]: print(df.dtypes)

A      int64
B      object
C     float64
dtype: object
```



PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Let's practice!



PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Training and Test Sets

Splitting up your dataset

```
In [1]: from sklearn.model_selection import train_test_split  
In [2]: X_train, X_test, y_train, y_test = train_test_split(X, y)
```

	X_train	y_train
0	1.0	n
1	4.0	n
2	7.0	n
3	2.0	n
4	5.0	n
5	5.0	n
6	6.0	n

	X_test	y_test
0	9.0	y
1	1.0	n
2	4.0	n

Stratified sampling

- 100 samples, 80 class 1 and 20 class 2
- Training set: 75 samples, 60 class 1 and 15 class 2
- Test set: 25 samples, 20 class 1 and 5 class 2

Stratified sampling

```
In [3]: y["labels"].value_counts()

class1    80
class2    20
Name: labels, dtype: int64

In [4]: X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=
In [5]: y_train["labels"].value_counts()

class1    60
class2    15
Name: labels, dtype: int64

In [6]: y_test["labels"].value_counts()

class1    20
class2     5
Name: labels, dtype: int64
```



PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Let's practice!